

Spotify Music Analysis

Team Fine Arts: Art and Music

12/7/22

Final Project Report

Authors

Kaushik Narasimha Bukkapatnam

Jonathan Olavarria

Katrina Arbogast

Varad Luktuke

Napathsawan Poopaiboon

Sudha Ramakrishnan

Introduction

Over the past few decades, there have been significant changes in the way the media arts are consumed. Everything from the projects that are produced, how they are produced, and the user experience has undergone an astonishing transformation in this new era of fast and limitless digital access. The media firms and individuals that have thrived and sustained are those that realized early on that usage trends and user attraction are the crucial elements to get a competitive edge over the competition.

Spotify is the industry leader in music streaming services, with over 100 million users and a market share of more than 31% among those who subscribe to online music. This motivated us to go further into the Spotify song database in order to uncover intriguing relationships between the songs and the artists as well as to gather insightful knowledge.

Problem Statement

If I were a struggling music producer and looking to produce my next music, what genres of music should I be targeting to produce based on the current market trends and are there any track features that have an impact the popularity of the track.

The approach

A two-fold approach is taken here :

- An overview of the genres and the artists
- An attempt to compute the track popularity based on its features

Data in scope

The data in scope is from the Spotify music platform that was extracted using spotify API. The Spotify Web API provides artist, album, and track data, as well as audio features and analysis, all easily accessible via the R package `spotiflyr`. Each music track has 23 audio attributes recorded, including descriptors like duration, tempo, key, and mode as well as confidence indicators like acousticness, liveness, speechiness, and instrumentality as well as perceptual indicators like energy, loudness, danceability, and valence (positiveness). The data used for this analysis is a dataset of over 30,000 songs obtained from GitHub. To learn more about the data, visit <https://github.com/rfordatascience/tidytuesday/tree/master/data/2020/2020-01-21>

Libraries and Data Input

```
suppressPackageStartupMessages(library(knitr))#Used to create a document that is a mixture  
suppressPackageStartupMessages(library(tidyverse))#for Data manipulation and Wrangling tasks  
suppressPackageStartupMessages(library(dplyr))#for Data manipulation and Wrangling tasks  
suppressPackageStartupMessages(library(ggplot2))#for effective visualizations  
suppressPackageStartupMessages(library(gridExtra))#for effective visualizations  
suppressPackageStartupMessages(library(reshape2))#used to melt the correlation matrix
```

Reading in the data

```
#Creating a "spotify_songs" datafram by reading the data
spotify_songs <- readr::read_csv('https://raw.githubusercontent.com/rfordatascience/tidytu

Rows: 32833 Columns: 23
-- Column specification -----
Delimiter: ","
chr (10): track_id, track_name, track_artist, track_album_id, track_album_na...
dbl (13): track_popularity, danceability, energy, key, loudness, mode, spec...
i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

spotify_songs

# A tibble: 32,833 x 23
  track_id      track~1 track~2 track~3 track~4 track~5 track~6 playl~7 playl~8
  <chr>        <chr>    <chr>    <dbl> <chr>    <chr>    <chr>    <chr>
  1 6f807x0ima9a~ I Don'~ Ed She~       66 2oCs0D~ I Don'~ 2019-0~ Pop Re~ 37i9dQ~
  2 0r7CVbZTWZgb~ Memori~ Maroon~       67 63rPSO~ Memori~ 2019-1~ Pop Re~ 37i9dQ~
  3 1z1Hg7Vb0AhH~ All th~ Zara L~       70 1HoSmj~ All th~ 2019-0~ Pop Re~ 37i9dQ~
  4 75FpbthrwQmz~ Call Y~ The Ch~       60 1nqYs0~ Call Y~ 2019-0~ Pop Re~ 37i9dQ~
  5 1e8PAfcKUYoK~ Someon~ Lewis ~       69 7m7vv9~ Someon~ 2019-0~ Pop Re~ 37i9dQ~
  6 7fvUMiyapMsR~ Beauti~ Ed She~       67 2yiy9c~ Beauti~ 2019-0~ Pop Re~ 37i9dQ~
  7 20AylPUDDfwR~ Never ~ Katy P~       62 7INHYS~ Never ~ 2019-0~ Pop Re~ 37i9dQ~
  8 6b1RNvAcJjQH~ Post M~ Sam Fe~       69 6703SR~ Post M~ 2019-0~ Pop Re~ 37i9dQ~
  9 7bF6tC03gFb8~ Tough ~ Avicii       68 7CvAfG~ Tough ~ 2019-0~ Pop Re~ 37i9dQ~
  10 1IXGILkPm0t0~ If I C~ Shawn ~      67 4Qxzbf~ If I C~ 2019-0~ Pop Re~ 37i9dQ~
# ... with 32,823 more rows, 14 more variables: playlist_genre <chr>,
#   playlist_subgenre <chr>, danceability <dbl>, energy <dbl>, key <dbl>,
#   loudness <dbl>, mode <dbl>, speechiness <dbl>, acousticness <dbl>,
#   instrumentalness <dbl>, liveness <dbl>, valence <dbl>, tempo <dbl>,
#   duration_ms <dbl>, and abbreviated variable names 1: track_name,
#   2: track_artist, 3: track_popularity, 4: track_album_id,
#   5: track_album_name, 6: track_album_release_date, 7: playlist_name, ...
# i Use `print(n = ...)` to see more rows, and `colnames()` to see all variable names

#getting the dimensions of the dataset
dim(spotify_songs)
```

```
[1] 32833    23
```

```
#take a quick look  
summary(spotify_songs)
```

track_id	track_name	track_artist	track_popularity
Length:32833	Length:32833	Length:32833	Min. : 0.00
Class :character	Class :character	Class :character	1st Qu.: 24.00
Mode :character	Mode :character	Mode :character	Median : 45.00
			Mean : 42.48
			3rd Qu.: 62.00
			Max. :100.00
track_album_id	track_album_name	track_album_release_date	
Length:32833	Length:32833	Length:32833	
Class :character	Class :character	Class :character	
Mode :character	Mode :character	Mode :character	
playlist_name	playlist_id	playlist_genre	playlist_subgenre
Length:32833	Length:32833	Length:32833	Length:32833
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character
danceability	energy	key	loudness
Min. :0.0000	Min. :0.000175	Min. : 0.000	Min. :-46.448
1st Qu.:0.5630	1st Qu.:0.581000	1st Qu.: 2.000	1st Qu.: -8.171
Median :0.6720	Median :0.721000	Median : 6.000	Median : -6.166
Mean :0.6548	Mean :0.698619	Mean : 5.374	Mean : -6.720
3rd Qu.:0.7610	3rd Qu.:0.840000	3rd Qu.: 9.000	3rd Qu.: -4.645
Max. :0.9830	Max. :1.000000	Max. :11.000	Max. : 1.275
mode	speechiness	acousticness	instrumentalness
Min. :0.0000	Min. :0.0000	Min. :0.0000	Min. :0.0000000
1st Qu.:0.0000	1st Qu.:0.0410	1st Qu.:0.0151	1st Qu.:0.0000000
Median :1.0000	Median :0.0625	Median :0.0804	Median :0.0000161
Mean :0.5657	Mean :0.1071	Mean :0.1753	Mean :0.0847472
3rd Qu.:1.0000	3rd Qu.:0.1320	3rd Qu.:0.2550	3rd Qu.:0.0048300
Max. :1.0000	Max. :0.9180	Max. :0.9940	Max. :0.9940000
liveness	valence	tempo	duration_ms

```

Min.    :0.0000  Min.    :0.0000  Min.    : 0.00  Min.    : 4000
1st Qu.:0.0927  1st Qu.:0.3310  1st Qu.: 99.96  1st Qu.:187819
Median   :0.1270  Median   :0.5120  Median   :121.98  Median   :216000
Mean     :0.1902  Mean     :0.5106  Mean     :120.88  Mean     :225800
3rd Qu.:0.2480  3rd Qu.:0.6930  3rd Qu.:133.92  3rd Qu.:253585
Max.    :0.9960  Max.    :0.9910  Max.    :239.44  Max.    :517810

```

Observation: There are a total of 32833 rows and 23 columns in the data

Lets look at what each attribute denotes-

```
#Creating a dataframe to hold the information about the attributes of the data
url <- 'https://raw.githubusercontent.com/vpcincin/DataWrangling/main/Data_Dictionary.csv'
spotify_attributes <- readr::read_csv(url)
```

```

Rows: 23 Columns: 3
-- Column specification -----
Delimiter: ","
chr (3): variable, class, description

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

print(spotify_attributes, n = 23)

# A tibble: 23 x 3
  variable            class      description
  <chr>              <chr>      <chr>
  1 track_id          character "Song unique ID"
  2 track_name         character "Song Name"
  3 track_artist        character "Song Artist"
  4 track_popularity    double     "Song Popularity (0-100) where higher is ~"
  5 track_album_id      character "Album unique ID"
  6 track_album_name    character "Song album name"
  7 track_album_release_date character "Date when album released"
  8 playlist_name       character "Name of playlist"
  9 playlist_id          character "Playlist ID"
 10 playlist_genre       character "Playlist genre"
 11 playlist_subgenre    character "Playlist subgenre"

```

12	danceability	double	"Danceability describes how suitable a tr~
13	energy	double	"Energy is a measure from 0.0 to 1.0 and ~
14	key	double	"The estimated overall key of the track. ~
15	loudness	double	"The overall loudness of a track in decib~
16	mode	double	"Mode indicates the modality (major or mi~
17	speechiness	double	"Speechiness detects the presence of spok~
18	acousticness	double	"A confidence measure from 0.0 to 1.0 of ~
19	instrumentalness	double	"Predicts whether a track contains no voc~
20	liveness	double	"Detects the presence of an audience in t~
21	valence	double	"A measure from 0.0 to 1.0 describing the~
22	tempo	double	"The overall estimated tempo of a track i~
23	duration_ms	double	"Duration of song in milliseconds"

We get the description of each variable and its class for understanding our data

Checking for null values

```
#getting the count of total null values in data
sum(is.na(spotify_songs))
```

```
[1] 15
```

```
#checking the null values by columns
colSums(is.na(spotify_songs))
```

	track_id	track_name	track_artist
	0	5	5
track_popularity		track_album_id	track_album_name
0		0	5
track_album_release_date		playlist_name	playlist_id
0		0	0
playlist_genre		playlist_subgenre	danceability
0		0	0
energy		key	loudness
0		0	0
mode		speechiness	acousticness
0		0	0
instrumentalness		liveness	valence
0		0	0
tempo		duration_ms	
0		0	

Observation: We see a total of 15 null values in the dataset. When checked deeper on the column level, we see that there are a total of 5 missing values for each of the columns- ‘track_name’, ‘track_artist’, ‘track_album_name’.

Checking for variable types

```
# checking variable types for consistencies
str(spotify_songs[])
```

```
tibble [32,833 x 23] (S3: tbl_df/tbl/data.frame)
$ track_id          : chr [1:32833] "6f807x0ima9a1j3VPbc7VN" "0r7CVbZTWZgbTCYdfa2P31"
$ track_name        : chr [1:32833] "I Don't Care (with Justin Bieber) - Loud Luxury"
$ track_artist      : chr [1:32833] "Ed Sheeran" "Maroon 5" "Zara Larsson" "The Chainsmokers"
$ track_popularity   : num [1:32833] 66 67 70 60 69 67 62 69 68 67 ...
$ track_album_id     : chr [1:32833] "2oCs0DGTsR098Gh5ZS12Cx" "63rPS0264uRjW1X5E6cWv6"
$ track_album_name    : chr [1:32833] "I Don't Care (with Justin Bieber) [Loud Luxury Remix]"
$ track_album_release_date: chr [1:32833] "2019-06-14" "2019-12-13" "2019-07-05" "2019-07-13"
$ playlist_name       : chr [1:32833] "Pop Remix" "Pop Remix" "Pop Remix" "Pop Remix" ...
$ playlist_id         : chr [1:32833] "37i9dQZF1DXcZDD7cfEKhW" "37i9dQZF1DXcZDD7cfEKhW"
$ playlist_genre       : chr [1:32833] "pop" "pop" "pop" "pop" ...
$ playlist_subgenre    : chr [1:32833] "dance pop" "dance pop" "dance pop" "dance pop" ...
$ danceability        : num [1:32833] 0.748 0.726 0.675 0.718 0.65 0.675 0.449 0.542 0.571 ...
$ energy              : num [1:32833] 0.916 0.815 0.931 0.93 0.833 0.919 0.856 0.903 0.911 ...
$ key                 : num [1:32833] 6 11 1 7 1 8 5 4 8 2 ...
$ loudness            : num [1:32833] -2.63 -4.97 -3.43 -3.78 -4.67 ...
$ mode                : num [1:32833] 1 1 0 1 1 1 0 0 1 1 ...
$ speechiness         : num [1:32833] 0.0583 0.0373 0.0742 0.102 0.0359 0.127 0.0623 0.091 ...
$ acousticness        : num [1:32833] 0.102 0.0724 0.0794 0.0287 0.0803 0.0799 0.187 0.196 ...
$ instrumentalness    : num [1:32833] 0.00 4.21e-03 2.33e-05 9.43e-06 0.00 0.00 0.00 4.82e-05 ...
$ liveness             : num [1:32833] 0.0653 0.357 0.11 0.204 0.0833 0.143 0.176 0.111 0.149 ...
$ valence              : num [1:32833] 0.518 0.693 0.613 0.277 0.725 0.585 0.152 0.367 0.311 ...
$ tempo                : num [1:32833] 122 100 124 122 124 ...
$ duration_ms          : num [1:32833] 194754 162600 176616 169093 189052 ...
```

Observations: The attribute ‘mode’ currently has a numeric field, however it is supposed to be a Boolean/factor variable, as it has values{0,1} The attribute ‘track_album_release_date’ is currently a character column but it is supposed to be a field with date values

It is important for us to change the type of these variables as they may be important for our analysis.

Cleaning the data (dealing with null values and modifying datatypes)

We previously inspected a total of 5 missing values each for the ‘track_name’, ‘track_artist’ and ‘track_album_name’ attributes. Let us impute these missing values by a constant value ‘unknown’ as these parameters would not be impacting our analysis since we are not focusing on these three attributes and moreover it is a very small fraction of the dataset.

Also, we will be changing the datatypes of the two attributes to the required type that we observed to be incorrect

```
#Missing Value Treatment
spotify_songs$track_artist[is.na(spotify_songs$track_artist)] <- 'unknown'
spotify_songs$track_album_name[is.na(spotify_songs$track_album_name)] <- 'unknown'
spotify_songs$track_name[is.na(spotify_songs$track_name)] <- 'unknown'
#check
sum(is.na(spotify_songs))

[1] 0

#Modifying Data types
spotify_songs$mode <- as.factor(spotify_songs$mode)
spotify_songs$track_album_release_date <- as.Date(spotify_songs$track_album_release_date)
#check
class(spotify_songs$mode)

[1] "factor"

class(spotify_songs$track_album_release_date)

[1] "Date"
```

Exploratory Data Analysis

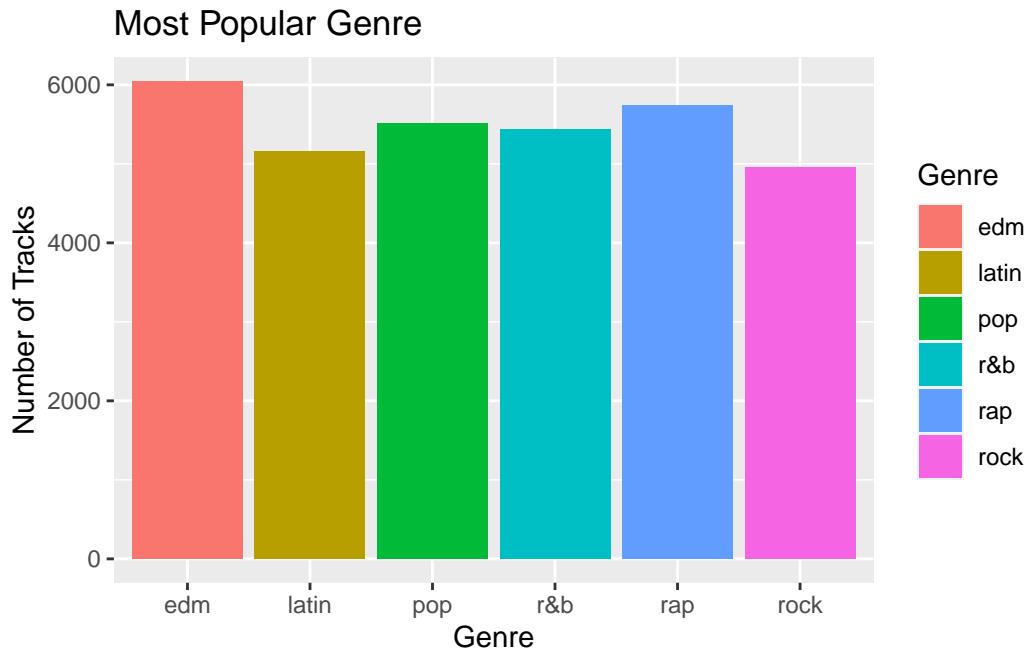
Firstly, we start by looking at the number of tracks in each genre to find out which genre has most number of releases.

```
spotify_songs %>% group_by(Genre = playlist_genre) %>%
  summarise(No_of_tracks = n()) %>% arrange(desc(No_of_tracks)) %>% knitr::kable()
```

Genre	No_of_tracks
edm	6043
rap	5746
pop	5507
r&b	5431
latin	5155
rock	4951

```
spotify_genres <- spotify_songs %>% group_by(Genre = playlist_genre) %>%
  summarise(No_of_tracks = n()) %>% arrange(desc(No_of_tracks)) %>%
  ggplot(aes(x = Genre, y = No_of_tracks, fill = Genre)) +
  geom_bar(stat = "identity") +
  labs(title = "Most Popular Genre", x = "Genre", y = "Number of Tracks")

spotify_genres
```

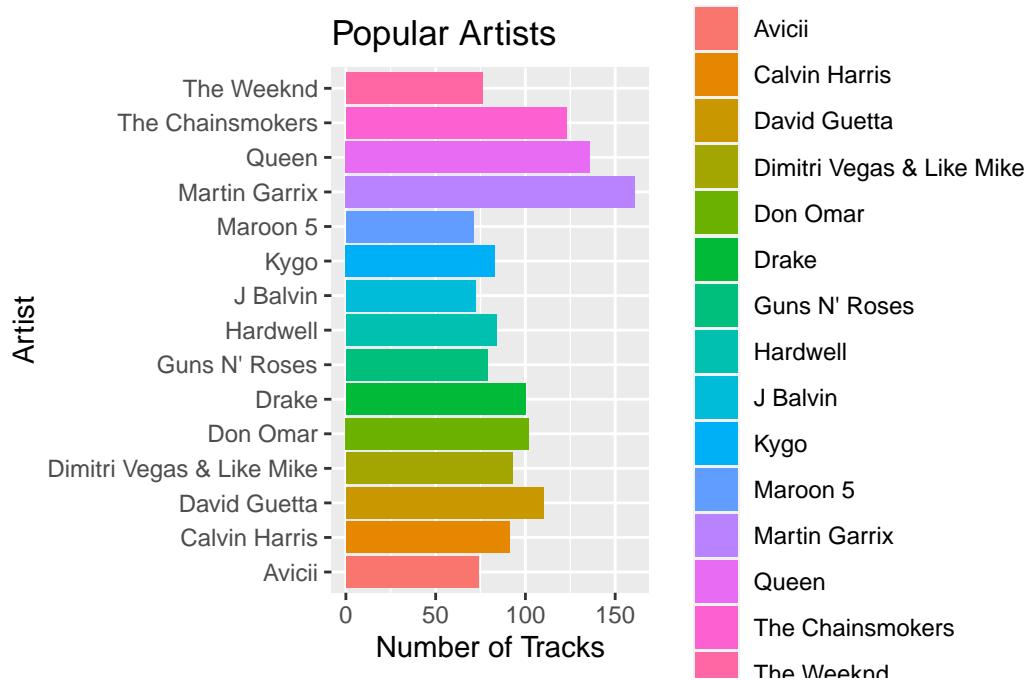


Observation: EDM songs quantify the most part of the data followed by rap and then pop depicting the listening trends for the respective genres. This indicates that most of the users listen to EDM music followed by Rap music.

Then we look at the number of songs released by each artist:

```
# artists with most releases
most_releases <- spotify_songs %>% group_by(Artist = track_artist) %>%
  summarise(No_of_tracks = n()) %>%
  arrange(desc(No_of_tracks)) %>%
  top_n(15, wt = No_of_tracks) %>%
  ggplot(aes(x = Artist, y = No_of_tracks, fill = Artist)) +
  geom_bar(stat = "identity") +
  coord_flip() +
  labs(title = "Popular Artists", x = "Artist", y = "Number of Tracks")
```

most_releases



Observation: With more than 150 songs Martin Garrix has been the most listened artist in this span followed by the band Queen with around 130 songs.

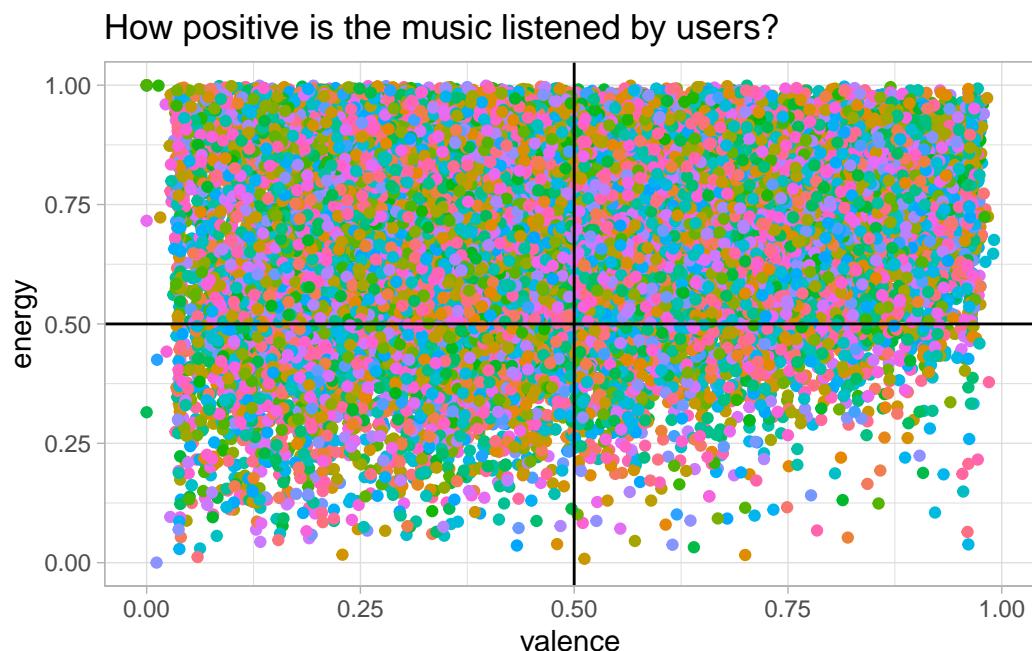
One of the main parameters of a song would be valence. It is a parameter to indicate musical positiveness conveyed by a track. Tracks with high valence sound more positive (e.g. happy, cheerful, euphoric), while tracks with low valence sound more negative (e.g. sad, depressed, angry). We look into the general trend of musical positiveness of tracks by plotting them against energy.

```

valence <- spotify_songs %>%
  ggplot(aes(x= valence, y= energy, color= track_name)) +
  geom_jitter(show.legend = FALSE) +
  geom_vline(xintercept = 0.5) +
  geom_hline(yintercept = 0.5) +
  scale_x_continuous(breaks= seq(0, 1, 0.25)) +
  scale_y_continuous(breaks= seq(0, 1, 0.25)) +
  labs(title= "How positive is the music listened by users?") +
  theme_light()

```

valence



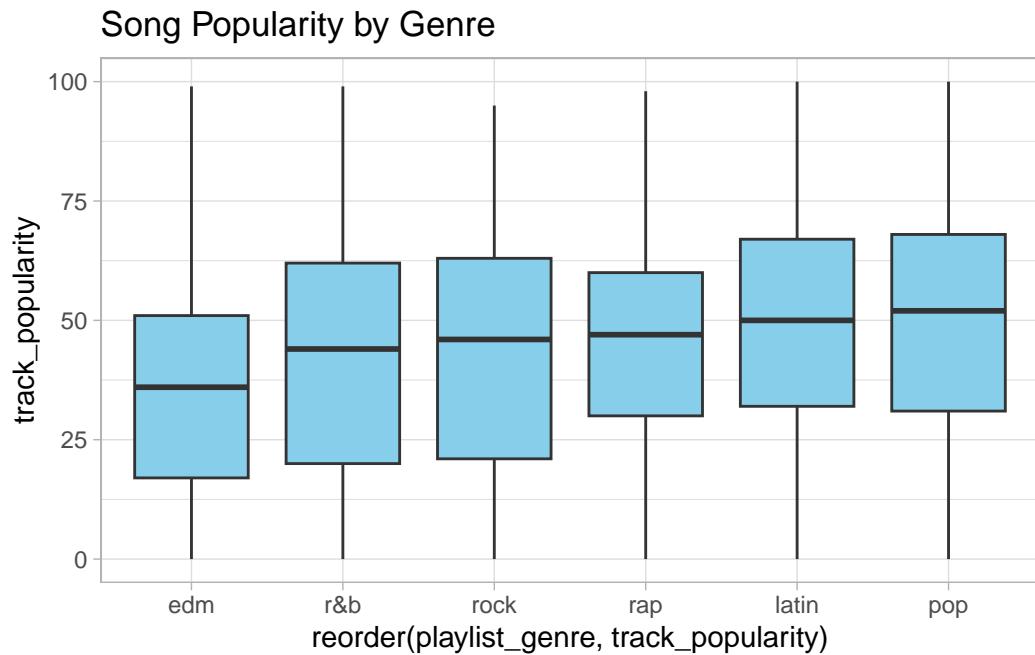
Observation: As seen from the plot, most of the songs are on the left side, which means they are sad, depressed, or angry. However, many of those tracks have high energy which can be interpreted as fast, loud, and noisy. Yet, there are still many tracks on the upper right, which represents the happy and positive.

Then we compared the popularity of songs by grouping them by genres to find out which genre would be more likely to gain popularity:

```

Pop_acc_genre <-spotify_songs%>%
  ggplot(aes(x=reorder(playlist_genre,track_popularity),y=track_popularity))+
  geom_boxplot(fill="skyblue")+
  ggtitle("Song Popularity by Genre")+
  theme_light()
Pop_acc_genre

```



Observation: The average song popularity index of pop songs is higher compared to the rest of the genres. Latin songs seconds the list.

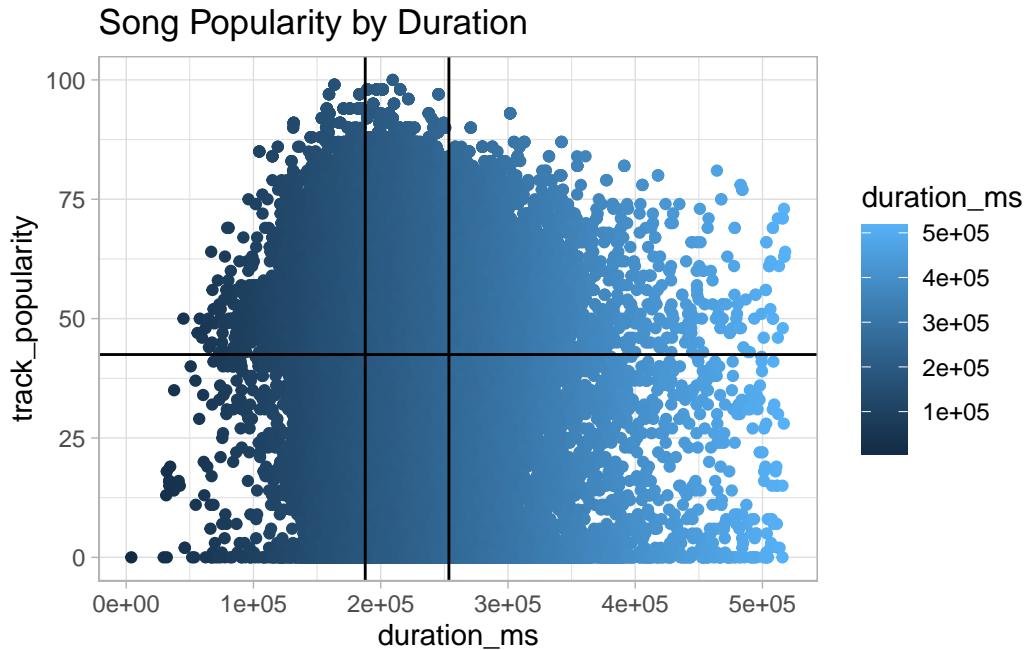
Next, we checked how duration of songs affected the popularity:

```

Popularity_duration <-spotify_songs%>%
  ggplot(aes(x=duration_ms,
             y=track_popularity, color = duration_ms))+ 
  geom_point()+
  geom_vline(xintercept = quantile(spotify_songs$duration_ms, 0.25)) +
  geom_vline(xintercept = quantile(spotify_songs$duration_ms, 0.75)) +
  geom_hline(yintercept = mean(spotify_songs$track_popularity)) +
  ggtitle("Song Popularity by Duration")+
  theme_light()

```

Popularity_duration



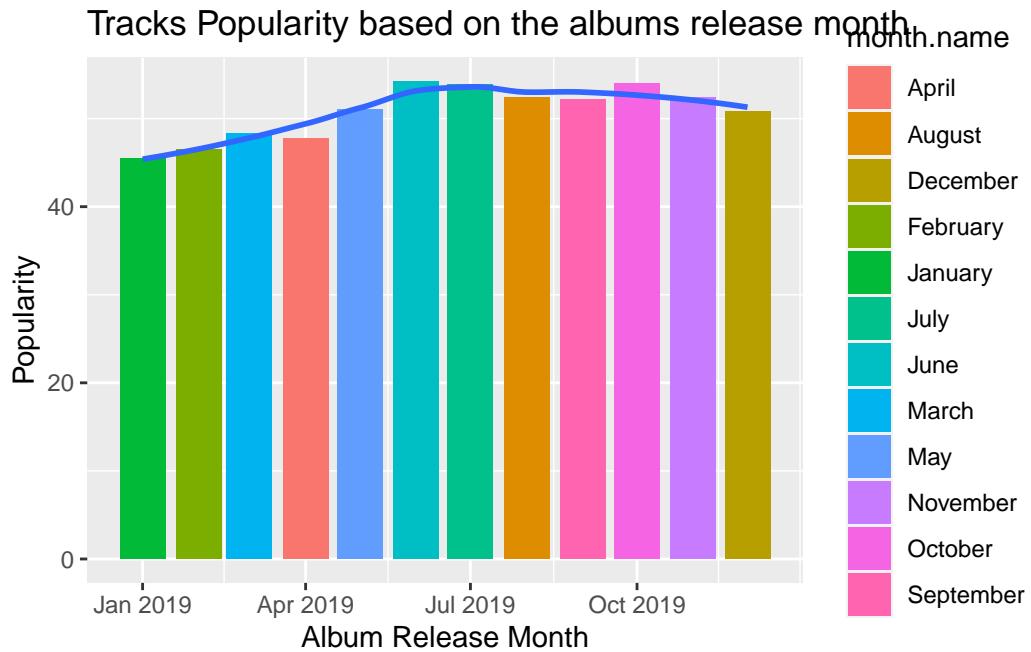
Observation: Songs that are in the upper 50% of popularity tend to localize around the center of song duration. We can see that between 187819 ms and 253585 ms most songs tend to accumulate.

Then we checked the correlation between track popularity and release date to find out whether the popularity of song is getting affected by the month it is released in:

```
top_dates_to_release <- spotify_songs %>%
  group_by(month = lubridate::floor_date(track_album_release_date, "month")) %>%
  summarise(avg_pop = mean(track_popularity))%>%
  filter(month >= '2019-01-01' & month < '2020-01-01')%>%
  ggplot(aes(x = month, y = avg_pop, fill = month.name)) +
  geom_bar(stat = "identity") +
  geom_smooth(se = FALSE, fill="blue")+
  labs(title = "Tracks Popularity based on the albums release month",
       x = 'Album Release Month',
       y = 'Popularity')

top_dates_to_release
```

```
`geom_smooth()` using method = 'loess' and formula = 'y ~ x'
```



Observation: The average track popularity for a given release month (in the year 2019), shows higher averages in months June through December compared to that of months January through May.

Checking the outliers

```
duration_boxplot <- spotify_songs %>%
  ggplot() +
  geom_boxplot(aes(duration_ms))

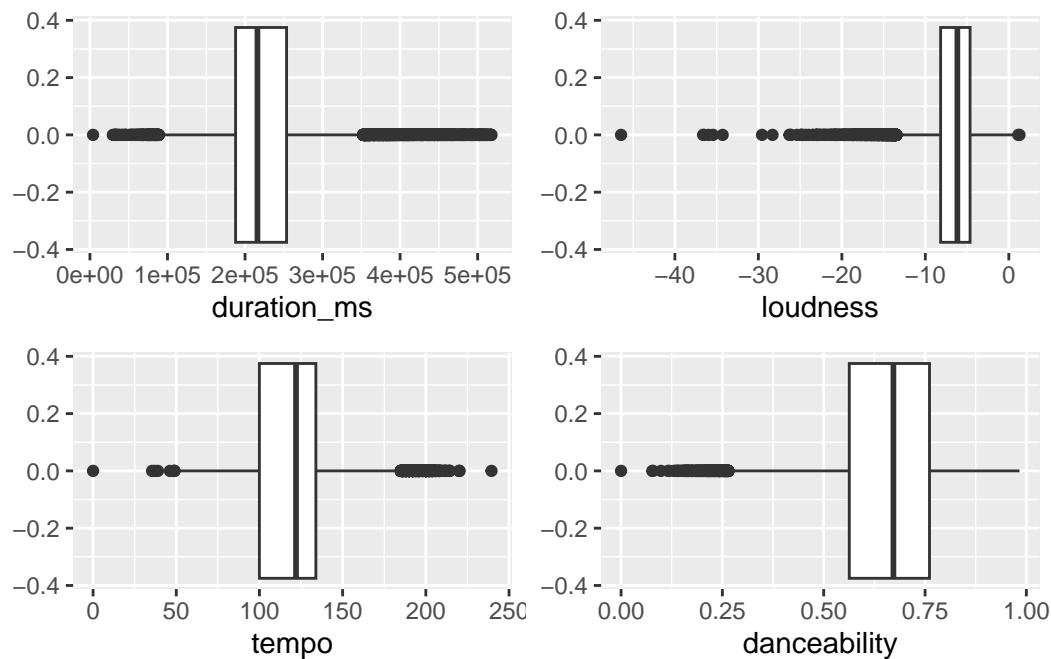
loudness_boxplot <- spotify_songs %>%
  ggplot() +
  geom_boxplot(aes(loudness))

tempo_boxplot <- spotify_songs %>%
  ggplot() +
  geom_boxplot(aes(tempo))

danceability_boxplot <- spotify_songs %>%
```

```
ggplot() +
  geom_boxplot(aes(danceability))
```

```
grid.arrange(duration_boxplot, loudness_boxplot, tempo_boxplot, danceability_boxplot, nrow=2)
```



```
dq1 <- quantile(spotify_songs$duration_ms, 0.25)
dq3 <- quantile(spotify_songs$duration_ms, 0.75)

lq1 <- quantile(spotify_songs$loudness, 0.25)
lq3 <- quantile(spotify_songs$loudness, 0.75)

tq1 <- quantile(spotify_songs$tempo, 0.25)
tq3 <- quantile(spotify_songs$tempo, 0.75)

danceq1 <- quantile(spotify_songs$danceability, 0.25)
danceq3 <- quantile(spotify_songs$danceability, 0.75)
```

Observation: For analysis we choose to look at these four variables for possible outliers, each of them indeed do have outliers. Duration, loudness, and tempo have outliers at both extremes, while danceability only has lower bound outliers.

The First and Third Quartiles for the variables are:

Quantile	Duration (ms)	Loudness	Tempo	Danceability
25%	187819	-8.171	99.96	0.563
75%	253585	-4.645	133.918	0.761

Plotting correlation heatmap

We then proceeded to check the correlations between our key attributes in order to examine if any attribute in general has more effect on the track popularity

```
#In order to check for correlations, we need to first segregate the numerical attributes
spotify_songs_num <- spotify_songs[,c(4,12,13,15,17,18,19,20,21,22,23)]
summary(spotify_songs_num)

track_popularity  danceability          energy          loudness
Min. : 0.00      Min. :0.0000      Min. :0.000175    Min. :-46.448
1st Qu.: 24.00    1st Qu.:0.5630    1st Qu.:0.581000  1st Qu.: -8.171
Median : 45.00    Median :0.6720    Median :0.721000  Median : -6.166
Mean   : 42.48    Mean   :0.6548    Mean   :0.698619  Mean   : -6.720
3rd Qu.: 62.00    3rd Qu.:0.7610    3rd Qu.:0.840000  3rd Qu.: -4.645
Max.   :100.00    Max.   :0.9830    Max.   :1.000000  Max.   :  1.275
speechiness       acousticness      instrumentalness  liveness
Min. :0.0000      Min. :0.0000      Min. :0.0000000  Min. :0.0000
1st Qu.:0.0410    1st Qu.:0.0151    1st Qu.:0.0000000  1st Qu.:0.0927
Median :0.0625    Median :0.0804    Median :0.0000161  Median :0.1270
Mean   :0.1071    Mean   :0.1753    Mean   :0.0847472  Mean   :0.1902
3rd Qu.:0.1320    3rd Qu.:0.2550    3rd Qu.:0.0048300  3rd Qu.:0.2480
Max.   :0.9180    Max.   :0.9940    Max.   :0.9940000  Max.   :0.9960
valence           tempo            duration_ms
Min. :0.0000      Min. : 0.00      Min. : 4000
1st Qu.:0.3310    1st Qu.: 99.96    1st Qu.:187819
Median :0.5120    Median :121.98    Median :216000
Mean   :0.5106    Mean   :120.88    Mean   :225800
3rd Qu.:0.6930    3rd Qu.:133.92    3rd Qu.:253585
Max.   :0.9910    Max.   :239.44    Max.   :517810

head(spotify_songs_num, 5)

# A tibble: 5 x 11
  track_p~1  dance~2  energy  loudn~3  speec~4  acous~5  instr~6  liven~7  valence  tempo
  <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>
```

```

<dbl> <dbl>
1     66  0.748  0.916 -2.63  0.0583  0.102  0      0.0653  0.518  122.
2     67  0.726  0.815 -4.97  0.0373  0.0724 4.21e-3  0.357  0.693  100.
3     70  0.675  0.931 -3.43  0.0742  0.0794 2.33e-5  0.11   0.613  124.
4     60  0.718  0.93   -3.78  0.102   0.0287 9.43e-6  0.204  0.277  122.
5     69  0.65   0.833 -4.67  0.0359  0.0803 0       0.0833  0.725  124.
# ... with 1 more variable: duration_ms <dbl>, and abbreviated variable names
#   1: track_popularity, 2: danceability, 3: loudness, 4: speechiness,
#   5: acousticness, 6: instrumentalness, 7: liveness
# i Use `colnames()` to see all variable names

```

```

#creating a correlation matrix restricting upto two decimal places
cormat <- round(cor(spotify_songs_num),2)
head(cormat)

```

	track_popularity	danceability	energy	loudness	speechiness
track_popularity	1.00	0.06	-0.11	0.06	0.01
danceability	0.06	1.00	-0.09	0.03	0.18
energy	-0.11	-0.09	1.00	0.68	-0.03
loudness	0.06	0.03	0.68	1.00	0.01
speechiness	0.01	0.18	-0.03	0.01	1.00
acousticness	0.09	-0.02	-0.54	-0.36	0.03
	acousticness	instrumentalness	liveness	valence	tempo
track_popularity	0.09	-0.15	-0.05	0.03	-0.01
danceability	-0.02	-0.01	-0.12	0.33	-0.18
energy	-0.54	0.03	0.16	0.15	0.15
loudness	-0.36	-0.15	0.08	0.05	0.09
speechiness	0.03	-0.10	0.06	0.06	0.04
acousticness	1.00	-0.01	-0.08	-0.02	-0.11
	duration_ms				
track_popularity	-0.14				
danceability	-0.10				
energy	0.01				
loudness	-0.12				
speechiness	-0.09				
acousticness	-0.08				

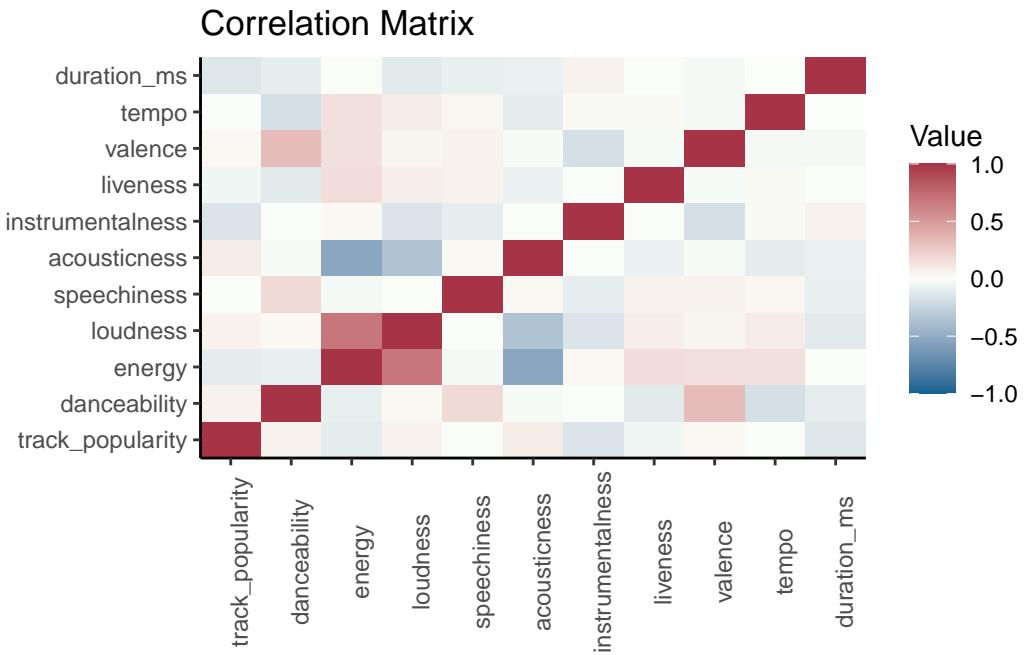
```

#melting the correlation matrix for plotting purpose
melted_cormat <- melt(cormat)
head(melted_cormat)

```

	Var1	Var2	value
1	track_popularity	track_popularity	1.00
2	danceability	track_popularity	0.06
3	energy	track_popularity	-0.11
4	loudness	track_popularity	0.06
5	speechiness	track_popularity	0.01
6	acousticness	track_popularity	0.09

```
#plotting the correlation matrix by using gradient fill
melted_cormat %>%
  ggplot(aes(Var1, Var2, fill=value)) +
  geom_tile() +
  labs(x = NULL, y = NULL, fill = "Value", title="Correlation Matrix") + scale_fill_gradient
#geom_text() +
theme_classic() +
scale_x_discrete(expand=c(0,0)) +
scale_y_discrete(expand=c(0,0)) +
theme(axis.text.x = element_text(angle = 90))
```



Observation: From the correlation matrix we can see that there is no single attribute/feature of a track having a high correlation with the popularity of the track. There is a very moderate correlation between the track popularity and the acousticness of the track but since the

correlation is weak, we cannot have the acoustics as a measure for determining the popularity of a track.

Linear Regression Model

Creating dummy variable for Playlist Genre

```
spotify_songs$genre_rap <- ifelse(spotify_songs$playlist_genre == 'rap', 1, 0)
spotify_songs$genre_rnb <- ifelse(spotify_songs$playlist_genre == 'r&b', 1, 0)
spotify_songs$genre_pop <- ifelse(spotify_songs$playlist_genre == 'pop', 1, 0)
spotify_songs$genre_latin <- ifelse(spotify_songs$playlist_genre == 'latin', 1, 0)
spotify_songs$genre_edm <- ifelse(spotify_songs$playlist_genre == 'edm', 1, 0)
```

Build Linear Regression Model

Y is track popularity.

Xs are danceability, energy, loudness, speechiness, acousticness, instrumentalness, liveness, valence, tempo, duration_ms, genre_rap, genre_rnb, genre_pop, genre_latin, and genre_edm

```
model <- lm(track_popularity ~ danceability + energy + loudness + speechiness + acousticne
#view regression model output
summary(model)
```

Call:

```
lm(formula = track_popularity ~ danceability + energy + loudness +
    speechiness + acousticness + instrumentalness + liveness +
    valence + tempo + duration_ms + genre_rap + genre_rnb + genre_pop +
    genre_latin + genre_edm, data = spotify_songs)
```

Residuals:

Min	1Q	Median	3Q	Max
-61.433	-17.261	3.199	18.627	67.271

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	7.993e+01	1.682e+00	47.533	< 2e-16 ***
danceability	1.051e+01	1.137e+00	9.240	< 2e-16 ***

```

energy           -2.938e+01  1.224e+00 -24.009  < 2e-16 ***
loudness        1.626e+00  6.479e-02  25.088  < 2e-16 ***
speechiness     -2.161e+00  1.477e+00 -1.463  0.143358
acousticness    2.894e+00  7.303e-01  3.962  7.44e-05 ***
instrumentalness -9.095e+00 6.463e-01 -14.072  < 2e-16 ***
liveness         -3.289e+00 8.744e-01 -3.761  0.000169 ***
valence          -8.947e-01 6.658e-01 -1.344  0.178992
tempo             2.371e-02 5.074e-03  4.673  2.98e-06 ***
duration_ms      -4.564e-05 2.284e-06 -19.981  < 2e-16 ***
genre_rap        -5.138e+00 5.623e-01 -9.138  < 2e-16 ***
genre_rnb        -6.649e+00 5.194e-01 -12.800  < 2e-16 ***
genre_pop         2.803e-01 4.998e-01  0.561  0.574914
genre_latin       -1.166e+00 5.307e-01 -2.197  0.028045 *
genre_edm        -9.408e+00 5.277e-01 -17.827  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 23.83 on 32817 degrees of freedom
 Multiple R-squared: 0.09053, Adjusted R-squared: 0.09012
 F-statistic: 217.8 on 15 and 32817 DF, p-value: < 2.2e-16

Observation: A linear regression model describes the relationship between a *dependent variable*, y , and one or more *independent variables*, X . With 32817 degrees of freedom, we found that the f value is 217.8 and p value was found to be less than 2.2e-16. The regression model has an extremely low adjusted R-squared value suggesting that the input variables are not good indicators of song popularity. This makes sense when referencing the correlation heat map as we saw very weak correlations between all the variables and song popularity.

Conclusion

We would like to start off by acknowledging some potential sources of bias in our analysis. First of which is that our data only covers songs from six genres, it is common knowledge that there are far more than six main genres in music and therefore we must preface our conclusion's with the fact that they only apply to songs in the genres of pop, edm, latin, R&B, rock, and rap. We cannot extrapolate these findings to songs in other genres. As far as our personal biases for the research, our team went in to this project thinking that the trends and correlations within the music would be very strongly defined. However, as the project matured and our analysis became more defined we realized that the trends may not be as significant as we initially thought they would be.

In our exploratory analysis we investigated the distribution of our data as well as interactions between some of the variables. The data was relatively uniformly distributed accross the 6

genres. However, we did find that at the artist level we had a few artists with far more songs in the data than some of the others these were Martin Garrix, Queen, and The Chainsmokers. This means that across the playlists sampled for to create our data set these artists were the most popular. In addition to this we looked at the interactions between variables. We found that there is a “sweet spot” for song duration meaning that songs that lie within 187819 ms and 253585 ms tend to be more popular than songs that lie outside of this range. We also found that songs that came from albums in June and July are considerably more popular than songs that came from albums released earlier in the year. Lastly we created a correlation heat map to see how our variables are correlated with each other and this is where discovered that there are fairly weak correlations between most of the variables. There are no strong correlations between any of the song features and popularity however, there are some strong correlations between a few of the features themselves. For example energy is positively correlated with loudness and acousticness is negatively correlated with both energy and loudness. These correlations are not so surprising because as music listeners as we know that energetic songs tend to be louder and acoustic songs are generally more quiet and have much less energy.

Moving forward with our analysis we fit a regression model to see if we could reliably predict song popularity based off of the various features in our data set. The multiple regression model had an adjusted R-squared value of 0.09 which is extremely low and means that we were only able to account for 9% of the variability in song popularity using song features as predictors. Due to this we can conclude that a song’s features are not a good indicator of the song’s popularity. This result is expected as our correlation heat map shows that there are no strong correlations between song features and popularity. To wrap things up through our exploratory analysis we were able to find interesting trends in the data relating to release date, genre, and artist popularity however, when it came to predicting song popularity there were not strong enough correlations in the data and therefore our model was fairly inaccurate.

```
sessionInfo()
```

```
R version 4.2.1 (2022-06-23)
```

```
Platform: x86_64-apple-darwin17.0 (64-bit)
```

```
Running under: macOS Big Sur ... 10.16
```

```
Matrix products: default
```

```
BLAS: /Library/Frameworks/R.framework/Versions/4.2/Resources/lib/libRblas.0.dylib
```

```
LAPACK: /Library/Frameworks/R.framework/Versions/4.2/Resources/lib/libRlapack.dylib
```

```
locale:
```

```
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
```

```
attached base packages:
```

```
[1] stats      graphics   grDevices utils      datasets  methods    base
```

```

other attached packages:
[1] reshape2_1.4.4   gridExtra_2.3   forcats_0.5.1   stringr_1.4.0
[5] dplyr_1.0.10    purrr_0.3.4     readr_2.1.2    tidyverse_1.3.2 knitr_1.41
[9] tibble_3.1.8    ggplot2_3.4.0   tidyverse_1.3.2 knitr_1.41

loaded via a namespace (and not attached):
[1] Rcpp_1.0.9          lattice_0.20-45   lubridate_1.8.0
[4] assertthat_0.2.1   digest_0.6.29    utf8_1.2.2
[7] R6_2.5.1           cellranger_1.1.0  plyr_1.8.7
[10] backports_1.4.1   reprex_2.0.1   evaluate_0.15
[13] highr_0.9          httr_1.4.3     pillar_1.8.0
[16] rlang_1.0.6         curl_4.3.2     googlesheets4_1.0.0
[19] readxl_1.4.0       rstudioapi_0.13 Matrix_1.4-1
[22] rmarkdown_2.14      splines_4.2.1   labeling_0.4.2
[25] googledrive_2.0.0  bit_4.0.4     munsell_0.5.0
[28] broom_1.0.0        compiler_4.2.1  modelr_0.1.8
[31] xfun_0.35          pkgconfig_2.0.3  mgcv_1.8-40
[34] htmltools_0.5.3    tidyselect_1.1.2 fansi_1.0.3
[37] crayon_1.5.1       tzdb_0.3.0     dbplyr_2.2.1
[40] withr_2.5.0        grid_4.2.1    nlme_3.1-157
[43] jsonlite_1.8.0     gtable_0.3.0   lifecycle_1.0.3
[46] DBI_1.1.3          magrittr_2.0.3  scales_1.2.0
[49] cli_3.4.1          stringi_1.7.8 vroom_1.5.7
[52] farver_2.1.1       fs_1.5.2      xml2_1.3.3
[55] ellipsis_0.3.2     generics_0.1.3 vctrs_0.5.1
[58] tools_4.2.1         bit64_4.0.5   glue_1.6.2
[61] hms_1.1.1          parallel_4.2.1 fastmap_1.1.0
[64] yaml_2.3.5         colorspace_2.0-3 gargle_1.2.0
[67] rvest_1.0.2         haven_2.5.0

```